

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/326972628>

Applying machine learning techniques to baseball pitch prediction

Conference Paper · January 2014

DOI: 10.5220/0004763905200527

CITATIONS

14

READS

2,591

7 authors, including:



Phuong Hoang
CareerBuilder

9 PUBLICATIONS 104 CITATIONS

[SEE PROFILE](#)



Hien T. Tran
North Carolina State University

155 PUBLICATIONS 3,578 CITATIONS

[SEE PROFILE](#)

Applying Machine Learning Techniques to Baseball Pitch Prediction

Michael Hamilton¹, Phuong Hoang², Lori Layne³, Joseph Murray², David Padgett³, Corey Stafford⁴
and Hien Tran²

¹*Mathematics Department, Rutgers University, New Brunswick, New Jersey, U.S.A.*

²*Department of Mathematics, North Carolina State University, Raleigh, North Carolina, U.S.A.*

³*MIT Lincoln Laboratory, Lexington, Massachusetts, U.S.A.*

⁴*Department of Applied Physics & Applied Mathematics, Columbia University, New York, New York, U.S.A.*

Keywords: Pitch Prediction, Feature Selection, ROC, Hypothesis Testing, Machine Learning.

Abstract: Major League Baseball, a professional baseball league in the US and Canada, is one of the most popular sports leagues in North America. Partially because of its popularity and the wide availability of data from games, baseball has become the subject of significant statistical and mathematical analysis. Pitch analysis is especially useful for helping a team better understand the pitch behavior it may face during a game, allowing the team to develop a corresponding batting strategy to combat the predicted pitch behavior. We apply several common machine learning classification methods to PITCH f/x data to classify pitches by type. We then extend the classification task to prediction by utilizing features only known before a pitch is thrown. By performing significant feature analysis and introducing a novel approach for feature selection, moderate improvement over former results is achieved.

1 INTRODUCTION

Baseball is one of the most popular sports in North America. In 2012, Major League Baseball (MLB) had the highest season attendance of any American sports league (MLB, 2012). Partially due to this popularity and the discrete nature of gameplay (allowing easy recording of game statistics between plays) and the long history of baseball data collection, baseball has become the target of significant mathematical and statistical analysis. Player performance, for example, is often analyzed so baseball teams can modify their roster (by drafting and trading players) to achieve the best possible team configuration.

One area of statistical analysis of baseball that has gained attention in the last decade is pitch analysis. To aid this study, baseball pitch data produced by the PITCH f/x system is now widely available for both public and private use. This data contains useful information about each pitch; several characteristics such as pitch speed, break angle, and type are recorded. Because of the accessibility of large volumes of data, both fans and professionals can perform their own pitch studies, including sabermetrics analysis. Related pitch data analysis is available in the

literature. For example, Weinstein-Gould (Weinstein-Gould, 2009) examines pitching strategy of major league pitchers, specifically determining whether or not pitchers (from a game theoretic approach) implement optimally mixed strategies for handling batters. The research suggests that pitchers do mix optimally with respect to the *pitch* variable. In an economic sense, this means that pitchers behave rationally relative to the result of any given pitch. An interesting note from the author is that although MLB pitchers are in the perfect position to utilize optimal strategy mixing (compared to other research subjects who have little motivation to optimally mix strategies), "...experience, large monetary incentives, and a competitive environment are not sufficient conditions to compel players to play optimally." Knowing this, we obtain useful information about how pitchers make decisions and (theoretically) which factors are more important than others in prediction; by knowing the results pitchers respond to (and the ones they don't respond to), it is possible reverse engineer this information and correspondingly tweak predictions for maximal accuracy.

Count (number of balls and strikes in the current at bat) is often cited as a basis for decisive strategy.

For example, Ganeshapillai and Guttag (Ganeshapillai and Guttag, 2012) show that pitchers are much more predictable in counts that favor the batter (usually more balls than strikes). Furthermore, Hopkins and Magel (Hopkins and Magel, 2008) show a distinct effect of count on the slugging percentage of the batter. More specifically, they show that average slugging percentage is significantly lower in counts that favor the pitcher; however, there is no significant difference in average slugging percentage (a weighted measure of the on-base frequency of a batter) in neutral counts or counts that favor the batter (Hopkins and Magel, 2008). These results verify that count has a significant effect on the pitcher-batter relationship, and will thus be an important factor in pitch prediction. Another interesting topic is pitch prediction, which could have significant real-world applications and potentially improve batter performance in baseball. One example of research on this topic is the work by Ganeshapillai and Guttag (Ganeshapillai and Guttag, 2012), who use a linear support vector machine (SVM) to perform binary (*fastball* vs. *nonfastball*) classification on pitches of unknown type. The SVM is trained on PITCH f/x data from pitches in 2008 and tested on data from 2009. Across all pitchers, an average prediction accuracy of roughly 70 percent is obtained, though pitcher-specific accuracies vary.

In this paper we provide a machine learning approach to pitch prediction, using classification methods to predict pitch types. Our results build upon the work in (Ganeshapillai and Guttag, 2012); however we are able to improve performance by examining different types of classification methods and by taking a pitcher adaptive approach to feature set selection. For more information about baseball itself, consult the Appendix for a glossary of baseball terms.

2 METHODS

2.1 PITCH f/x Data

Our classifiers are trained and tested using PITCH f/x data from all MLB games during the 2008 and 2009 seasons. Raw data is publicly available (Pitchf/x, 2013), though we use scraping methods to transform the data into a suitable format. The data contains approximately 50 features (each represents some characteristic of a pitch like speed or position); however, we only use 18 features from the raw data and create additional features that are relevant to prediction. For example, some created features are: the percentage of fastballs thrown in the previous inning, the velocity of the previous pitch, strike result percentage of previous

pitch, and current game count (score). For a full list of features used, see Appendix.

We apply classification methods to the data to predict pitches. On that note, it is important to clarify a subtle distinction between pitch classification and pitch prediction. The distinction is simply that classification uses post-pitch information about a pitch to determine which type it is, whereas prediction uses pre-pitch information to classify its type. For example, we may use features like pitch speed and curve angle to determine whether or not it was a fastball. These features are not available pre-pitch; in that case we use information about prior results from the same scenario to judge which pitch can be expected.

The prediction process is performed as binary classification (see section 2.2); all pitch types are members of one of two classes (*fastball* and *nonfastball*). We conduct prediction for all pitchers who had at least 750 pitches in both 2008 and 2009. This specification results in a set of 236 pitchers. For each pitcher, the data is further split by each count and, with 12 count possibilities producing 2,832 smaller subsets of data (one for each pitcher and count combination). After performing feature selection (see section 2.3) on each data subset, each classifier (see Appendix) is trained on each subset of data from 2008 and tested on each subset of data from 2009. The average classification accuracy for each classifier is computed for test points with a type confidence (one feature in the PITCH f/x data that measures the confidence level that the pitch type is correct) of at least 0.5.

2.2 Classification Methods

Classification is the process of taking an unlabeled data observation and using some rule or decision-making process to assign a label to it. Within the scope of this research, classification represents determining the type of a pitch, i.e. given a pitch x with characteristics x_i , determine which pitch type (curveball, fastball, slider, etc) x is. There are several classification methodologies one can use to accomplish this task, here we used the methods of Support Vector Machine (SVM) and k -nearest neighbor (k -NN), they are explained in full detail in (Theodoridis and Koutroumbas, 2009).

2.3 Feature Selection Methods

The key differences between our approach and former research (Ganeshapillai and Guttag, 2012) is the feature selection methodology. Rather than using a static set of features, a different optimal set of features is

used for each pitcher/count pair. This allows the algorithm to adapt for optimal performance on each individual subset of data.

In baseball there are a number of factors that influence the pitcher's decision (consciously or unconsciously). For example, one pitcher may not like to throw curveballs during the daytime because the increased visibility makes them easier to spot; however, another pitcher may not make his pitching decisions based on the time of the game. In order to maximize accuracy of a prediction model, one must try to accommodate each of these factors. For example, a pitcher may have particularly good control of a certain pitch and thus favors that pitch, but how can one create a feature to represent its favorability? One could, for example, create a feature that measures the pitcher's success with a pitch since the beginning of the season, or the previous game, or even the previous batter faced. Which features would best capture the true effect of his preference for that pitch? The answer is that each of these approaches may be best in different situations, so they all must be considered for maximal accuracy. Pitchers have different dominant pitches, strategies and experiences; in order to maximize accuracy our model must be adaptable to various pitching situations.

Of course, simply adding many features to our model is not necessarily the best choice because we leave noise from situationally useless features and suffer from curse of dimensionality issues. We change our problem of predicting a pitch into predicting a pitch for each given pitcher in each given count. We maximize accuracy by choosing (for each pitcher/count pair) an optimal pool of features from the entire available set. This allows us to maintain our adaptive strategy while controlling dimensionality.

2.3.1 Feature Selection Implementation

As alluded to earlier, our feature selection scheme is adaptive, finding a good feature set for each (pitcher, count), e.g., (Rivera, 0-2). Our adaptive feature selection approach consists of 3 stages.

1. We create about 80 features (see the full list in the Appendix). We then group all of our generated features together into groups by similarity. Such grouping might contain 4 to 12 similar features such as accuracy over the last pitch, the last five pitches, ten pitches, and accuracy over all pitches in the last inning, etc.
2. We then compute the Receiver Operating Characteristic (ROC) curve for each group of features, then select the strongest one or two to move on to the next stage. In practice, selecting only the best

feature provides worse prediction than selecting the best two or three features. Hence at this stage, the size of each group is reduced from 4 to 12 to 2 or 3.

3. We next remove all redundant features from our final feature set. From our grouping, features are taken based on their relative strength. There is the possibility that a group of features might not have good predictive power. In those cases, the resulting set of features is pruned by conducting hypothesis testing to measure significance of each feature at the $\alpha = .01$ level.

For a schematic diagram of our approach, see Figure 1.

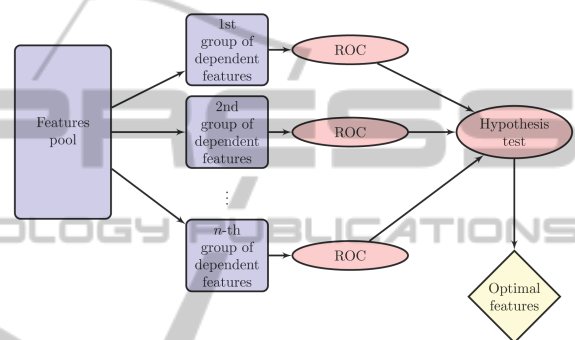


Figure 1: Schematic diagram of the proposed features selection.

2.3.2 ROC Curves

The Receiver Operating Characteristic (ROC) curve are used for each individual feature in order to measure how useful a feature is in prediction. We calculate this by measuring the area between the single feature ROC curve and the line created by standard guessing. This value of area tells us how much better the feature is at distinguishing the two classes, compared to standard guessing. These area values are in the range of $[0, 0.5]$, where a value of 0 represents no improvement over random guessing and 0.5 would represent perfect distinction between both classes. For a more detailed description of the ROC curve, see (Fawcett, 2006) and also Figure 2.

2.3.3 Hypothesis Testing

The ability of a feature to distinguish between two classes can be determined using a hypothesis test. Given any feature f , we compare μ_1 and μ_2 , the mean values of f in Class 1 (*fastballs*) and Class 2 (*nonfastballs*), respectively. Then we consider

$$H_0 : \mu_1 = \mu_2 \quad (1)$$

$$H_A : \mu_1 \neq \mu_2 \quad (2)$$




Figure 2: The diagonal line represents the tradeoff between true positive rate (sensitivity) and false positive rate by random guessing. The curve represents a shift in this tradeoff by using a given feature to assign class labels instead of randomly guessing. In this case the shift represents an improvement in distinction between the two classes and the region between the curve and line quantifies this improvement.

and conduct a hypothesis test using the student's t distribution. We compare the p -value of the test against a significance level of $\alpha = .01$. When the p -value is less than α , we reject the null hypothesis and conclude that the studied feature means are different for each class, meaning that the feature is significant in separating the classes. In that sense, this test allows us to remove features which have insignificant separation power.

3 RESULTS

In this paper, we propose a new technique in the problem of baseball pitch prediction. Specifically, we segment the prediction task by pitcher and count because each of these situations is different enough that it would be a mistake to consider them equally. For prediction by pitcher, we used data from 236 pitchers (as noted in section 2.1). We then selected eight pitchers from the 2008 and 2009 MLB regular seasons to examine in details.

Table 1: Data for each pitcher.

Pitcher	Training Size	Test Size
Fuentes	919	798
Madson	975	958
Meche	2821	1822
Park	1309	1178
Rivera	797	850
Vaquez	2412	2721
Wakefield	2110	1573
Weathers	943	813

Table 1 describes the training and testing sets. Data from 2008 season were used for training and data from 2009 were used for testing. Table 2 depicts the prediction accuracy among the eight pitchers as compared across classifiers as well as naive guess. On average, 79.76% of pitches are correctly predicted by SVM-L and SVM-G classifiers while k -NN perform slightly better, at 80.88% accurate. Furthermore, k -NN is also a better choice in term of computational speed, as noted in Table 3.

We compare the results of our prediction model to a naive model that predicts simply by guessing a pitcher's most common pitch (either *fastball* or *non-fastball*) from the training data and compute the improvement in accuracy in Table 4. The improvement factor (percent), I , is calculated as follow:

$$I = \frac{A_1 - A_0}{A_0} \times 100 \quad (3)$$

where A_0 and A_1 denotes the accuracies of naive guess and our model accordingly. The naive guess simply return the most frequent pitch type thrown by each pitcher, calculated from the training set (Ganeshapillai and Guttag, 2012).

The average prediction accuracy of our model over all 236 pitchers in 2009 season is 77.45%. Compared to the naive model's natural prediction accuracy, our model on average achieves 20.85% improvement. In previous work (Ganeshapillai and Guttag, 2012), the average prediction accuracy of 2009 season is 70% with 18% improvement over the naive model. It should be noted that previous work considers 359 pitchers who threw at least 300 pitches in both 2008 and 2009 seasons. In our study, we only consider those pitchers that threw at least 750 pitches in a season. This reduces the number of pitchers that we considered to 236 pitchers.

In addition, we demonstrate the prediction accuracy of our model for each count situation. As illustrated in Figures 3 and 4, prediction accuracy is significant higher in batter-favored counts and is approximately equal in neutral and pitcher-favored counts.

We also calculate prediction accuracy for 2012 season, using training data from 2011 season or from both 2010 and 2011 seasons. We again only select pitchers who had at least 750 pitches in those seasons. As shown in Table 5, the average pitch prediction accuracy is about 75% in both cases (even though the size of training data is double in the second case).

Table 2: Prediction accuracy comparison (percents). Symbols: k -Nearest Neighbors (k -NN), Support Vector Machine with linear kernel (SVM-L), Support Vector Machine with Gaussian kernel (SVM-G), Naive Guess (NG).

Classifier	Fuentes	Madson	Meche	Park	Rivera	Vaquez	Wakefield	Weathers	Average
k -NN	80.15	81.85	72.73	70.31	93.51	72.50	100.00	76.01	80.88
SVM-L	78.38	77.23	74.83	72.40	89.44	72.50	95.50	77.76	79.76
SVM-G	76.74	79.38	74.17	71.88	90.14	73.05	96.33	76.38	79.76
NG	71.05	25.56	50.77	52.60	89.63	51.20	100.00	35.55	66.04

Table 3: CPU Times (seconds).

Classifier	Fuentes	Madson	Meche	Park	Rivera	Vaquez	Wakefield	Weathers	Average
k -NN	0.3459	0.3479	0.3927	0.3566	0.4245	0.4137	0.4060	0.3480	0.3794
SVM-L	0.7357	0.5840	1.2616	0.7322	0.6441	1.1282	0.3057	0.5315	0.7408
SVM-G	0.3952	0.4076	0.7270	0.4591	0.4594	0.7248	0.5267	0.3641	0.5799




Figure 3: Prediction accuracy by count of 2009 season.




Figure 4: Prediction accuracy by count of 2012 season.

4 CONCLUSIONS AND FUTURE WORK

Originally our scheme developed from consideration of the factors that affect pitching decisions. For example, the pitcher/batter handedness matchup is often mentioned by sports experts as an effect, and it was originally included in our model. However, it was discovered that implementing segmentation of data based on handedness has essentially no effect on the

Table 4: Improvement over Naive Guess (percents)

Pitcher	Improvement	Classifier
Fuentes	12.81	k -NN
Madson	22.01	k -NN
Meche	47.39	SVM-L
Park	37.62	SVM-L
Rivera	0.04	k -NN
Vaquez	42.68	SVM-G
Wakefield	0.00	k -NN
Weathers	118.73	SVM-L

prediction results. Thus, handedness is no longer implemented as a further splitting criterion of the model, but this component remains a considered feature. In general, unnecessary data segmentations have negative impact solely because it reduce the size of training and testing data for classifiers to work with.

Most notable is our method of feature selection which widely varies the set of features used in each situation. Features that yield strong prediction in some situations fail to provide any benefit in others. In fact, it is interesting to note that in the 2008/2009 prediction scheme, every feature is used in at least one situation and no feature is used in every situation.

It is also interesting to note that the most successful classification algorithm of this model is supported by our feature selection technique. In general, Bayesian classifiers rely on a feature independence assumption, which is realistically not satisfied. However, our model survives this assumption because although the features within each of the 6 groups are highly dependent across groups. Thus the features which are ultimately chosen are highly independent.

The model represents a significant improvement over simple guessing. It is a useful tool for batting coaches, batters, and others who wish to understand the potential pitching implications of a given game

Table 5: Prediction model results for additional years. Note that percentage improvement is calculated on a per-pitcher basis and then averaged overall.

Train Year(s)	Test Year	Naive Guessing Accuracy	Our Model Accuracy	Improvement
2011	2012	62.72	75.20	24.82
2010 and 2011	2012	62.97	75.27	24.07

scenario. For example, batters could theoretically use this model to increase their batting average, assuming that knowledge about a pitch's type makes it easier to hit. The model, for example, is especially useful in certain intense game scenarios and achieves accuracy as high as 90 percent. It is in these game environments that batters can most effectively use this model to translate knowledge into hits. Additionally, it is interesting to note that in 0-2 counts where naive guessing is least accurate, our model performs relatively well.

Looking forward, much can be done to improve the model. First, new features would be helpful. There is much game information that we did not include in our model, such as batting averages, slugging percentage per batter, stadium location, weather, and others, which could help improve the prediction accuracy of the model. One potential modification is extension to multi-class classification. Currently, our model makes a binary decision and decides if the next pitch will be a fastball or not. It does not determine what kind of fastball the pitch may be. However, this task is much more difficult and would almost certainly result in a decrease in accuracy. Further, prediction is not limited to only the pitch type. For example, one could consider the prediction problem of determining where the pitch will be thrown (for example, a specific quadrant). Or it may be possible to predict, if in the given situation a hit were to occur, where the ball is likely to land. That information could be useful to prepare the corresponding defensive player for the impending flight of the ball.

ACKNOWLEDGEMENTS

This research was conducted with support from NSA grant H98230-12-1-0299 and NSF grants DMS-1063010 and DMS-0943855. We would like to thank Jessica Gronsbell for her useful advice and Mr. Tom Tippet, Director of Baseball Information Services for the Boston Red Sox, for meeting with us to discuss this research and providing us with useful feedback. Also we would like to thank Minh Nhat Phan for helping us scrape the PITCH f/x data.

REFERENCES

- Fawcett, T. (2006). An introduction to ROC analysis. *Pattern recognition letters*, 27(8):861–874.
- Ganeshapillai, G. and Guttag, J. (2012). Predicting the next pitch. In *MIT Sloan Sports Analytics Conference*.
- Hopkins, T. and Magel, R. (2008). Slugging percentage in differing baseball couns. *Journal of Quantitative Analysis in Sports*, 4(2):1136.
- MATLAB (2013). MATLAB documentation.
- MLB (2012). Major league baseball attendance records. Retrieved June 19, 2013 from <http://espn.go.com/mlb/attendance/year/2012>.
- Pitchf/x (2013). MLB pitch f/x data. Retrieved July, 2013 from <http://www.mlb.com>.
- SVM (2013). Support vector machines explained. Retrieved July 3, 2013 from <http://www.tristanfletcher.co.uk>.
- Theodoridis, S. and Koutroumbas, K. (2009). *Pattern recognition, fourth edition*. Academic Press, Burlington, Mass., 4th edition.
- Weinstein-Gould, J. (2009). Keeping the hitter off balance: Mixed strategies in baseball. *Journal of Quantitative Analysis in Sports*, 5(2):1173.
- Wikipedia (2013). Wikipedia glossary of baseball. Retrieved July, 2013 from http://en.wikipedia.org/wiki/Glossary_of_baseball.

APPENDIX

Generic (Original) Features

The original 18 useful features from the raw data.

1. At-bat-number: number of pitches recorded against a specific batter.
2. Outs: number of outs during an at-bat.
3. Batter's I.D.
4. Pitcher's I.D.
5. Pitcher Handedness: pitching hand of pitcher, i.e R = Right, L = Left.
6. Pitch-event: outcome of one pitch from the pitcher's perspective (ball, strike, hit-by-pitch, foul, in-play, etc.)
7. Hitter-event: outcome of the at-bat from the batter's perspective (ground-out, double, single, walk, etc.
8. Outcome.

9. Pitch-type: classification of pitch type, i.e FF = Four-seam Fastball, SL = Slider, etc.
10. Time-and-date
11. Start-speed: pitch speed, in miles per hours, measured from the initial position.
12. x-position: horizontal location of the pitch as it crosses the home plate.
13. y-position: vertical location of the pitch as it crosses the home plate.
14. On-first: binary column; display 1 if runner on first, 0 otherwise.
15. On-second: binary column; display 1 if runner on third, 0 otherwise.
16. On-third: binary column; display 1 if runner on third, 0 otherwise.
17. Type-confidence: a rating corresponding to the likelihood of the pitch type classification.
18. Ball-strike: display either ball or strike (not always clear from pitch-event).
18. Time (day/afternoon/night)
19. Vertical position of previous pitch thrown
20. Previous pitch in specific count: vertical position
21. Number of outs
22. Previous pitch: ball or strike (boolean)
23. Previous pitch in specific count: ball/strike (boolean)
24. Percentage of fastballs thrown in last inning pitched by pitcher
25. Previous pitch: pitch type
26. Previous pitch in specific count: pitch type
27. Percentage of fastballs over previous 5 pitches thrown to specific batter
28. Strike result percentage (SRP) (a metric we created that measures the percentage of strikes from all pitches in the given situation) of fastballs thrown in the previous inning
29. Previous pitch: fastball of nonfastball (boolean)
30. Previous pitch in specific count: fast-ball/nonfastball (boolean)
31. Percentage of fastballs over previous 10 pitches thrown to specific batter
32. SRP of nonfastballs thrown in previous inning
33. Previous 2 pitches: average of velocities
34. Previous 2 pitches in specific count: velocity average
35. Percentage of fastballs over previous 15 pitches thrown to specific batter
36. Percentage of fastballs thrown in the previous game pitched by pitcher
37. Previous 2 pitches: average of horizontal positions
38. Previous 2 pitches in specific count: horizontal position average
39. Percentage of fastballs over previous 5 pitches thrown in specific count
40. SRP of fastballs thrown in previous game
41. Previous 2 pitches: average of vertical positions
42. Previous 2 pitches in specific count: vertical position average
43. Percentage of fastballs over previous 10 pitches thrown in specific count
44. SRP of nonfastballs thrown in previous game
45. Previous 2 pitches: ball/strike combo
46. Previous 2 pitches in specific count: ball strike and combo

Additional Features

From the original features above, we create the following features.

1. Inning
2. Lifetime percentage of fastballs thrown by pitcher
3. Previous 3 pitches: averages of horizontal and vertical positions
4. Previous 3 pitches in specific count: horizontal and vertical position averages
5. Player on first base, second base, and third base (boolean)
6. Percentage of fastballs historically thrown to batter
7. Percentage of fastballs thrown in batter's previous at bat
8. Numeric score of result from previous meeting of current pitcher and batter
9. Previous 3 pitches: fastball and nonfastball combo, ball and strike combo
10. Previous 3 pitches in specific count: fastball and nonfastball combo, ball and strike combo
11. Weighted base score
12. Velocity of previous pitch
13. Previous pitch in specific count: velocity
14. Percentage of fastballs over previous 5, 10, 15, and 20 pitches
15. Number of base runners
16. Horizontal position of previous pitch thrown
17. Previous pitch in specific count: horizontal position
18. Time (day/afternoon/night)
19. Vertical position of previous pitch thrown
20. Previous pitch in specific count: vertical position
21. Number of outs
22. Previous pitch: ball or strike (boolean)
23. Previous pitch in specific count: ball/strike (boolean)
24. Percentage of fastballs thrown in last inning pitched by pitcher
25. Previous pitch: pitch type
26. Previous pitch in specific count: pitch type
27. Percentage of fastballs over previous 5 pitches thrown to specific batter
28. Strike result percentage (SRP) (a metric we created that measures the percentage of strikes from all pitches in the given situation) of fastballs thrown in the previous inning
29. Previous pitch: fastball of nonfastball (boolean)
30. Previous pitch in specific count: fast-ball/nonfastball (boolean)
31. Percentage of fastballs over previous 10 pitches thrown to specific batter
32. SRP of nonfastballs thrown in previous inning
33. Previous 2 pitches: average of velocities
34. Previous 2 pitches in specific count: velocity average
35. Percentage of fastballs over previous 15 pitches thrown to specific batter
36. Percentage of fastballs thrown in the previous game pitched by pitcher
37. Previous 2 pitches: average of horizontal positions
38. Previous 2 pitches in specific count: horizontal position average
39. Percentage of fastballs over previous 5 pitches thrown in specific count
40. SRP of fastballs thrown in previous game
41. Previous 2 pitches: average of vertical positions
42. Previous 2 pitches in specific count: vertical position average
43. Percentage of fastballs over previous 10 pitches thrown in specific count
44. SRP of nonfastballs thrown in previous game
45. Previous 2 pitches: ball/strike combo
46. Previous 2 pitches in specific count: ball strike and combo

47. Percentage of fastballs over previous 15 pitches thrown in specific count
48. Percentage fastballs thrown in previous at bat
49. Previous 2 pitches: fastball/nonfastball combo
50. Previous 2 pitches in specific count: fastball and nonfastball combo
51. SRP of previous 5 fastballs to specific batter
52. Numeric score for last at bat event
53. Previous 3 pitches: average of velocities
54. Previous 3 pitches in specific count: velocity average
55. SRP of previous 5 nonfastballs to specific batter
56. Batter handedness (boolean)
57. Cartesian quadrant for previous pitch
58. Cartesian quadrant average for previous 2 pitches
59. Cartesian quadrant average for previous 3 pitches
60. Fastball SRP over previous 5, 10, and 15 pitches
61. Nonfastball SRP over previous 5, 10, and 15 pitches
62. Cartesian quadrant for previous pitch in specific count
63. Cartesian quadrant average for previous 2 pitches in specific count
64. Cartesian quadrant average for previous 3 pitches in specific count
3. Ball: When the batter does not swing at the pitch and the pitch is outside the strike zone. If the batter accrues four balls in an at bat he gets a walk, a free pass to first base.
4. Hit-By-Pitch: When the pitch hits the batters body. The batter gets a free pass to first base, similar to a walk.
5. Hit: When the batter makes contact with the pitch and successfully reaches first, second or third base. Types of hits include single (batter ends at first base), doubles (batter ends at second base), triple (batter ends at third base) and home-run.
6. Out: When a batter or base runner cannot, for whatever reason, advance to the next base. Examples include striking out (batter can not advance to first), grounding out, popping out and lining out.
7. Count: Is the number of balls and strikes during an at bat. There are 12 possible counts spanning every combination of 0-3 balls (4 balls is a walk) and 0-2 strikes (3 strikes is a strikeout).
8. Run: When a base runner crosses home plate. This is a point for that player's team. The outcome of a baseball game is determined by which team has more runs at the end of nine Innings.
9. Inning: One of nine periods of playtime in a standard game.
10. Slugging Percentage: A measure of hitter power. Defined as the average number of bases the hitter earns per at bat.
11. Batting Average: The percentage of time the batter earns a hit.
12. At Bat : A series of pitches thrown by the pitcher to one hitter resulting in a hit, a walk, or an out.

Baseball Glossary and Info

Most of these definitions are obtained directly from (Wikipedia, 2013).

1. Strike Zone: A box over the home plate which defines the boundaries through which a pitch must pass in order to count as a strike when the batter does not swing. A pitch that does not cross the plate through the strike zone is a ball.
2. Strike: When a batter swings at a pitch but fails to hit it, when a batter does not swing at a pitch that is thrown within the strike zone, when the ball is hit foul and the strike count is less than 2 (a batter cannot strike out on a foul ball, however he can fly out), when a ball is bunted foul, regardless of the strike count, when the ball touches the batter as he swings at it, when the ball touches the batter in the strike zone, or when the ball is a foul tip. Three strikes and the batter is said to have struck out.

Classification Methods

Here we list the classifiers used in the experiment. For more information about these classification methods, consult the cited reference.

1. *k*-NN(s): *k*-nearest-neighbors algorithm (MATLAB: *knnsearch*) with standardized Euclidean distance metric (MATLAB, 2013).
2. SVM-L: Support Vector Machine with linear kernel (SVM, 2013)
3. SVM-G: Support Vector Machine with RBF (Gaussian) kernel (SVM, 2013)